

# Data.Table Package

## Table of Contents

### Base R

- In base R, we use the 'dataframe' to organize data.
- We also use `read.csv()` to import data. This can be slow though...
- You may also be familiar with the tidyverse, or specifically `dplyr` for managing data in R.
- Why is base R slow?
  - Bad r code (not pre-allocating result vectors)
  - Lazy copying when manipulating dataframes (like renaming)
  - Slower sort algorithms
  - Not leveraging sort order

### The Data.Table Package

- A different approach to dataframes in R
- `data.table` replaces the `data.frame`
- A `data.table` object is also a `data.frame`
  - You can use a `data.table` anywhere you can use a `data.frame`
- Goals of the `data.table` package
  - Reduce time to write code through sensical syntax (opinion)
  - Reduce the time code takes to run through efficient copying ordered grouping and ordered joins (fact)

- Some of the computing time gains depend on sorting (setting keys) the data.table then leveraging the sort order for joins, group by calculations, and subsetting

## Creating a data.table

- We can directly create a data.table similar to a data.frame
- We can easily convert a data.frame to a data.table
- We can use fread() to read in structured text as a data.table, this replaces read.csv()

```
library(data.table)
```

```
## We can directly create a data.table just like a dataframe
```

```
dt1 <-
```

```
  data.table(x1=1:10,  
            x2=10:1)
```

```
dt1
```

```
##      x1 x2  
##  1:  1 10  
##  2:  2  9  
##  3:  3  8  
##  4:  4  7  
##  5:  5  6  
##  6:  6  5  
##  7:  7  4  
##  8:  8  3  
##  9:  9  2  
## 10: 10  1
```

```
str(dt1)
```

```
## Classes 'data.table' and 'data.frame':  10 obs. of  2 variables:  
## $ x1: int  1 2 3 4 5 6 7 8 9 10  
## $ x2: int 10 9 8 7 6 5 4 3 2 1  
## - attr(*, ".internal.selfref")=<externalptr>
```

```
is.data.frame(dt1)
```

```
## [1] TRUE
```

```
is.data.table(dt1)
```

```
## [1] TRUE
```

## Creating a data.table 2

- We can easily convert a data.frame to a data.table
- We can use fread to read in structured text as a data.table (this replaces read.csv)

```
library(data.table)
```

```
library(plyr) # for the baseball data
```

```
## We can convert existing dataframe to dt
```

```
data(baseball)
```

```
str(baseball)
```

```
## 'data.frame':    21699 obs. of  22 variables:
## $ id      : chr  "ansonca01" "forceda01" "mathebo01" "startjo01" ...
## $ year    : int   1871 1871 1871 1871 1871 1871 1871 1872 1872
1872 ...
## $ stint   : int    1 1 1 1 1 1 1 1 1 1 ...
## $ team    : chr   "RC1" "WS3" "FW1" "NY2" ...
## $ lg      : chr    "" "" "" "" ...
## $ g       : int   25 32 19 33 29 29 29 46 37 25 ...
## $ ab      : int  120 162 89 161 128 146 145 217 174 130 ...
## $ r       : int   29 45 15 35 35 40 36 60 26 40 ...
## $ h       : int   39 45 24 58 45 47 37 90 46 53 ...
## $ X2b     : int   11 9 3 5 3 6 5 10 3 11 ...
## $ X3b     : int    3 4 1 1 7 5 7 7 0 0 ...
## $ hr      : int    0 0 0 1 3 1 2 0 0 0 ...
## $ rbi     : int   16 29 10 34 23 21 23 50 15 16 ...
## $ sb      : int    6 8 2 4 3 2 2 6 0 2 ...
## $ cs      : int    2 0 1 2 1 2 2 6 1 2 ...
## $ bb      : int    2 4 2 3 1 4 9 16 1 1 ...
## $ so      : int    1 0 0 0 0 1 1 3 1 0 ...
## $ ibb     : int   NA NA NA NA NA NA NA NA NA NA ...
## $ hbp     : int   NA NA NA NA NA NA NA NA NA NA ...
## $ sh      : int   NA NA NA NA NA NA NA NA NA NA ...
## $ sf      : int   NA NA NA NA NA NA NA NA NA NA ...
## $ gidp    : int   NA NA NA NA NA NA NA NA NA NA ...
```

```
baseball_dt <-
```

```
  data.table(baseball)
```

```

## Fast File Read
cohort <- fread("patient_demo.csv")

## how much faster?
system.time(cohort <- read.csv("patient_demo.csv"))

##      user  system elapsed
## 0.147    0.012    0.161

library(readr)
system.time(cohort <- read_csv("patient_demo.csv"))

##
## — Column specification

```

---

```

## cols(
##   patient_id = col_double(),
##   hosp_id = col_double(),
##   age = col_double(),
##   gender = col_character(),
##   mortality = col_character(),
##   SES = col_character()
## )

##      user  system elapsed
## 0.226    0.004    0.334

system.time(cohort <- fread("patient_demo.csv"))

##      user  system elapsed
## 0.038    0.000    0.023

```

## data.table Commands

- Recall a data.table is also a data.frame so all the data.frame methods work
- Data.table prints only the first 5 and last 5 rows of a table if the table is larger than 100 rows

```

## list all data.tables and their size
tables()

```

```

##           NAME      NROW NCOL MB
COLS KEY
## 1: baseball_dt 21,699   22  2
id,year,stint,team,lg,g,...

```

```
## 2:      cohort 100,000      6  3
patient_id,hosp_id,age,gender,mortality,SES
## 3:      dt1      10      2  0
x1,x2
## Total: 5MB
```

```
## Confirm this is a data.table (and a data frame)
class(cohort)
```

```
## [1] "data.table" "data.frame"
```

```
## the usual methods for data.frames all work for
summary(cohort)
```

```
##      patient_id      hosp_id      age      gender
##  Min.   :100000    Min.    :  1.00    Min.    : 22.00   Length:100000
##  1st Qu.:125000    1st Qu.: 25.00    1st Qu.: 58.00
##  Class :character
##  Median :150000    Median : 51.00    Median : 65.00
##  Mode  :character
##  Mean   :150000    Mean    : 50.55    Mean    : 64.93
##  3rd Qu.:174999    3rd Qu.: 75.00    3rd Qu.: 72.00
##  Max.   :199999    Max.    :100.00    Max.    :109.00
##
##      mortality      SES
##  Length:100000    Length:100000
##  Class :character  Class :character
##  Mode  :character  Mode  :character
##
##
##
```

```
## Data.tables print the first and last 5 rows
cohort
```

```
##      patient_id hosp_id age gender mortality SES
##      1:      100000      77  70   Male      Yes  Med
##      2:      100001      45  64   Male      No   Low
##      3:      100002       5  76   Male      No   Med
##      4:      100003      13  51   Male      No  High
```

```
##      5:      100004      70  76  Male      No  Low
##      ---
## 99996:      199995      60  66  Male      No  Med
## 99997:      199996      14  71  Male      No  Low
## 99998:      199997      39  73  Male      Yes Med
## 99999:      199998      21  60 Female     No  High
## 100000:      199999      81  61  Male      No  Low
```

## The i= Argument

- We will use the square brackets to interact with data.tables (the subset operator)
- Forget everything you know about the square brackets (sorry!), it is different in data.table
- The first argument to the square brackets is i=, which we use to subset the rows of the data.table
- This can be an integer vector or a Boolean vector
- Any references to variables will be first resolved within the data.table environment

```
cohort[] # prints all
```

```
##      patient_id hosp_id age gender mortality  SES
##      1:      100000      77  70  Male      Yes  Med
##      2:      100001      45  64  Male      No   Low
##      3:      100002       5  76  Male      No   Med
##      4:      100003      13  51  Male      No  High
##      5:      100004      70  76  Male      No   Low
##      ---
## 99996:      199995      60  66  Male      No   Med
## 99997:      199996      14  71  Male      No   Low
## 99998:      199997      39  73  Male      Yes  Med
## 99999:      199998      21  60 Female     No  High
## 100000:      199999      81  61  Male      No   Low
```

```
## the first row
```

```
cohort[i=1]
```

```
##      patient_id hosp_id age gender mortality SES
## 1:      100000      77  70   Male          Yes Med
```

```
## we usually don't bother naming the argument with i=
## Rows 10 to 20
cohort[10:20]
```

```
##      patient_id hosp_id age gender mortality  SES
## 1:      100009      64  71   Male          No  Low
## 2:      100010      43  83   Male          No  Low
## 3:      100011      62  66   Male          No  Low
## 4:      100012      54  58   Male          No  Med
## 5:      100013      69  82   Male          No  Low
## 6:      100014      20  72   Male          No  Med
## 7:      100015      46  48   Male          No High
## 8:      100016      19  71   Male          No High
## 9:      100017       1  70   Male          No  Med
## 10:     100018      81  60   Male          No  Med
## 11:     100019       5  76   Male          No  Med
```

## The i= Argument

```
## all patients older than 50
cohort[age>50]
```

```
##      patient_id hosp_id age gender mortality  SES
## 1:      100000      77  70   Male          Yes Med
## 2:      100001      45  64   Male          No  Low
## 3:      100002       5  76   Male          No  Med
## 4:      100003      13  51   Male          No High
## 5:      100004      70  76   Male          No  Low
## ---
## 92508:     199995      60  66   Male          No  Med
## 92509:     199996      14  71   Male          No  Low
## 92510:     199997      39  73   Male          Yes Med
## 92511:     199998      21  60 Female          No High
## 92512:     199999      81  61   Male          No  Low
```

```
## all patients older than 50 and male
cohort[age>50 & gender == "Male"]
```

```
##      patient_id hosp_id age gender mortality  SES
## 1:      100000      77  70   Male          Yes Med
## 2:      100001      45  64   Male          No  Low
```





```
## 5: NA NA NA NA
## 6: NA NA NA NA
## 7: NA NA NA NA
## 8: NA NA NA NA
## 9: NA NA NA NA
## 10: NA NA NA NA
```

```
baseball_dt[year > 2000]
```

```
##           id year stint team lg   g  ab   r   h X2b X3b hr rbi
sb cs bb so
##      1: alomaro01 2001      1  CLE AL 157 575 113 193  34  12 20 100
30  6 80 71
##      2: alomasa02 2001      1  CHA AL  70 220  17  54   8   1  4  21
1  2 12 17
##      3: anderbr01 2001      1  BAL AL 131 430  50  87  12   3  8  45
12  4 60 77
##      4: baineha01 2001      1  CHA AL  32  84   3  11   1   0  0   6
0  0  8 16
##      5: bellda01 2001      1  SEA AL 135 470  62 122  28   0 15  64
2  1 28 59
##      ---

## 1266: benitar01 2007      2  FLO NL  34   0   0   0   0   0  0   0
0  0  0  0
## 1267: benitar01 2007      1  SFN NL  19   0   0   0   0   0  0   0
0  0  0  0
## 1268: ausmubr01 2007      1  HOU NL 117 349  38  82  16   3  3  25
6  1 37 74
## 1269: aloumo01 2007      1  NYN NL  87 328  51 112  19   1 13  49
3  0 27 30
## 1270: alomasa02 2007      1  NYN NL   8  22   1   3   1   0  0   0
0  0  0  3
##           ibb hbp sh sf gidp
##      1:    5   4  9  9   9
##      2:    1   2  3  2   6
##      3:    4   8  2  1   3
##      4:    0   0  0  2   2
##      5:    1   3  5  4   8
##      ---
## 1266:    0   0  0  0   0
## 1267:    0   0  0  0   0
## 1268:    3   6  4  1  11
```

```
## 1269: 5 2 0 3 13
## 1270: 0 0 0 0 0
```

```
baseball_dt[team %in% c("NYN","CLE")]
```

```
##           id year stint team lg  g  ab  r   h X2b X3b hr rbi sb
## cs bb  so
##    1: hallmbi01 1901      1  CLE AL   5  19  2   4   0   0  0   3  0
## NA  2  NA
##    2: mooreea01 1901      1  CLE AL  31  99  5  16   0   0  0   6  1
## NA  6  NA
##    3: weyhigu01 1901      1  CLE AL   2   5  0   0   0   0  0   0  0
## NA  0  NA
##    4: hickmch01 1902      2  CLE AL 102 426 61 161  31  11  8  94  8
## NA 12  NA
##    5: lajoina01 1902      2  CLE AL  86 348 81 132  35   5  7  64 19
## NA 19  NA
##    ---
```

```
## 1461: easleda01 2007      1  NYN NL  76 193 24  54   6   0 10  26  0
## 1 19  35
## 1462: delgaca01 2007      1  NYN NL 139 538 71 139  30   0 24  87  4
## 0 52 118
## 1463: coninje01 2007      2  NYN NL  21  41  2   8   2   0  0   5  0
## 0  7   8
## 1464: aloumo01 2007      1  NYN NL  87 328 51 112  19   1 13  49  3
## 0 27  30
## 1465: alomasa02 2007      1  NYN NL   8  22  1   3   1   0  0   0  0
## 0  0   3
```

```
##           ibb hbp sh sf gidp
##    1:  NA    0  0 NA    NA
##    2:  NA    0  3 NA    NA
##    3:  NA    0  0 NA    NA
##    4:  NA    3  8 NA    NA
##    5:  NA    6  8 NA    NA
##    ---
```

```
## 1461: 1 5 0 1 2
## 1462: 8 11 0 6 12
## 1463: 2 0 1 1 1
## 1464: 5 2 0 3 13
## 1465: 0 0 0 0 0
```

## Data.Table: the “j” argument

- The second argument to the square brackets is the j argument
- We leave the first argument blank (the i argument) if we want all the rows and just put a comma
- The j argument can be any valid R expression
- If the j argument resolves to a list, a data.table is returned

```
## Returns age as a vector
```

```
cohort[,j=age][1:10]
```

```
## [1] 70 64 76 51 76 60 55 66 75 71
```

```
## we typically omit the j= piece
```

```
cohort[,age][1:10]
```

```
## [1] 70 64 76 51 76 60 55 66 75 71
```

```
## Adding list() returns a data.table
```

```
cohort[,list(age)]
```

```
##      age
```

```
## 1: 70
```

```
## 2: 64
```

```
## 3: 76
```

```
## 4: 51
```

```
## 5: 76
```

```
## ---
```

```
## 99996: 66
```

```
## 99997: 71
```

```
## 99998: 73
```

```
## 99999: 60
```

```
## 100000: 61
```

```
cohort[,list(age,patient_id)]
```

```
##      age patient_id
```

```
## 1: 70      100000
```

```
## 2: 64      100001
```

```
## 3: 76      100002
```

```
## 4: 51      100003
```

```
## 5: 76      100004
```

```
## ---
```

```
## 99996: 66      199995
```

```
## 99997: 71      199996
```

```
## 99998: 73      199997
```

```
## 99999: 60      199998
## 100000: 61      199999

## use . as a shortcut for list here
cohort[,.(age,patient_id)]

##      age patient_id
## 1:  70    100000
## 2:  64    100001
## 3:  76    100002
## 4:  51    100003
## 5:  76    100004
## ---
## 99996: 66      199995
## 99997: 71      199996
## 99998: 73      199997
## 99999: 60      199998
## 100000: 61      199999
```

## Data.Table: the “j” argument

Question: From the Baseball data, return a data.table with the rows from after the year 2000 with only the id, year and team columns using the list() command

Question: Repeat this using the . to replace the list command

Question: How many rows meet this condition?

- Answer

```
baseball_dt[year > 2000,list(id,year,team)]
```

```
##      id year team
## 1: alomaro01 2001  CLE
## 2: alomasa02 2001  CHA
## 3: anderbr01 2001  BAL
## 4: baineha01 2001  CHA
## 5: bellda01 2001  SEA
## ---
## 1266: benitar01 2007  FLO
## 1267: benitar01 2007  SFN
## 1268: ausmubr01 2007  HOU
## 1269: aloumo01 2007  NYN
## 1270: alomasa02 2007  NYN
```

```
baseball_dt[year > 2000,.(id,year,team)]
```

```
##           id year team
##    1: alomaro01 2001  CLE
##    2: alomasa02 2001  CHA
##    3: anderbr01 2001  BAL
##    4: baineha01 2001  CHA
##    5:  bellda01 2001  SEA
##    ---
## 1266: benitar01 2007  FLO
## 1267: benitar01 2007  SFN
## 1268: ausmubr01 2007  HOU
## 1269: aloumo01 2007  NYN
## 1270: alomasa02 2007  NYN

## How many rows?
nrow(baseball_dt[year > 2000,.(id,year,team)])

## [1] 1270
```

## Using Integers and Characters in the j argument

- In base R we can use integers or characters to reference columns as the second argument following [
- The same behavior is true in data.table (this was not always the case in prior versions of data.table)
- However, one difference is the class of the return object. Notice data.table returns a data.table, while the data.frame would return a vector instead of a dataframe.

```
# These work as expected
cohort[,1]
```

```
##           patient_id
##    1:         100000
##    2:         100001
##    3:         100002
##    4:         100003
##    5:         100004
##    ---
## 99996:        199995
## 99997:        199996
## 99998:        199997
```

```
## 99999:      199998
## 100000:     199999
```

```
cohort[, "age"]
```

```
##      age
## 1:  70
## 2:  64
## 3:  76
## 4:  51
## 5:  76
## ---
## 99996: 66
## 99997: 71
## 99998: 73
## 99999: 60
## 100000: 61
```

*# These do not!*

```
var_num <- 1
cohort[, var_num]
```

## Error in `[.data.table`(cohort, , var\_num): j (the 2nd argument inside [...]) is a single symbol but column name 'var\_num' is not found. Perhaps you intended DT[, ..var\_num]. This difference to data.frame is deliberate and explained in FAQ 1.1.

```
cohort[, ..var_num] # works
```

```
##      patient_id
## 1:    100000
## 2:    100001
## 3:    100002
## 4:    100003
## 5:    100004
## ---
## 99996: 199995
## 99997: 199996
## 99998: 199997
## 99999: 199998
## 100000: 199999
```

```
var_name <- "age"
cohort[, var_name]
```

```
## Error in `[.data.table`(cohort, , var_name): j (the 2nd argument
inside [...]) is a single symbol but column name 'var_name' is not
found. Perhaps you intended DT[, ..var_name]. This difference to
data.frame is deliberate and explained in FAQ 1.1.
```

```
cohort[,..var_name] # works
```

```
##           age
##      1:   70
##      2:   64
##      3:   76
##      4:   51
##      5:   76
##      ---
## 99996:   66
## 99997:   71
## 99998:   73
## 99999:   60
## 100000:  61
```

## More Complex j arguments

- As stated the j argument can be any valid R expression
- If we return a list of items, it will be coerced into a data.table
- We can name the items of the list, and it will be inherited to column names

```
## here is the average and SD of age
```

```
cohort[,mean(age)]
```

```
## [1] 64.93313
```

```
cohort[,sd(age)]
```

```
## [1] 10.00578
```

```
## Returned as a data.table
```

```
cohort[,list(mean(age),  
             sd(age))]
```

```
##           V1           V2
## 1: 64.93313 10.00578
```

```
## Name each column in the result
```

```
cohort[,list(avg=mean(age),  
             stdev=sd(age))]
```

```
##          avg      stdev
## 1: 64.93313 10.00578
```

## The j expression

Question: Calculate the average home runs (hr variable) and return the result as a data.table  
Question: Calculate the total number of at bats (ab variable) and the total number of rbi's (rbi variable) and the total number of players with a team value of "HOU". Name these columns in the resulting data.table

- Answer

```
baseball_dt[,list(mean(hr))]
```

```
##          V1
## 1: 5.234204
```

```
baseball_dt[,list(tot_ab = sum(ab),
                  tot_rbi = sum(rbi,na.rm=T),
                  tot_hous = sum(team=="HOU"))]
```

```
##      tot_ab tot_rbi tot_hous
## 1: 4891061  641644      497
```

## The j argument, complex expressions

- We can pass in more complex expressions in the j expression
- As before, if the result is not a simple list, it will not be a data.table
- This should remind you of the with() command in base r dataframes

```
## here is a table of gender by mortality
cohort[,table(gender,mortality)]
```

```
##          mortality
## gender      No   Yes
## Female  8399 1533
## Male   73612 16456
```

```
## With row proportions
cohort[,prop.table(table(gender,mortality),1)]
```



```
##           mortality
## gender           No           Yes
## Female 0.8456504 0.1543496
## Male   0.8172936 0.1827064

## chi-square test on gender and mortality
cohort[,chisq.test(gender,mortality)]

##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data: gender and mortality
## X-squared = 48.566, df = 1, p-value = 3.194e-12
```

## The j expression, reducing results

- As an example of a more complex argument to j lets perform a logistic regression on this dataset

```
## logistic regression on mortality
mod_result_1 <-
  cohort[,glm(mortality == "Yes" ~ age + gender + SES,
              family = "binomial")]
summary(mod_result_1)

##
## Call:
## glm(formula = mortality == "Yes" ~ age + gender + SES, family =
"binomial")
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.9126  -0.6536  -0.6099  -0.5495   2.1987
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.9637216  0.0640235 -46.291  < 2e-16 ***
## age          0.0189530  0.0008296  22.846  < 2e-16 ***
## genderMale   0.2065857  0.0291515   7.087 1.37e-12 ***
## SESLow       0.0305045  0.0232105   1.314   0.189
## SESMed       0.0144423  0.0223347   0.647   0.518
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 94245   on 99999   degrees of freedom
## Residual deviance: 93667   on 99995   degrees of freedom
## AIC: 93677
##
## Number of Fisher Scoring iterations: 4

## all in one step
cohort[,summary(glm(mortality == "Yes" ~ age + gender + SES,
                    family = "binomial"))]

##
## Call:
## glm(formula = mortality == "Yes" ~ age + gender + SES, family =
"binomial")
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.9126  -0.6536  -0.6099  -0.5495   2.1987
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.9637216  0.0640235 -46.291  < 2e-16 ***
## age          0.0189530  0.0008296  22.846  < 2e-16 ***
## genderMale   0.2065857  0.0291515   7.087 1.37e-12 ***
## SESLow       0.0305045  0.0232105   1.314   0.189
## SESMed       0.0144423  0.0223347   0.647   0.518
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 94245   on 99999   degrees of freedom
## Residual deviance: 93667   on 99995   degrees of freedom
## AIC: 93677
##
## Number of Fisher Scoring iterations: 4
```

## Using the `j` argument for variable transformations

- We may wish to have a transformed version of a variable

- We can name simply make a list of the variables we want, and add in the transformations

```
## Lets select patient_id, age, and add log age
```

```
cohort[,list(patient_id,
             age,
             l_age=log(age))]
```

```
##      patient_id age    l_age
## 1:      100000  70 4.248495
## 2:      100001  64 4.158883
## 3:      100002  76 4.330733
## 4:      100003  51 3.931826
## 5:      100004  76 4.330733
## ---
## 99996:      199995  66 4.189655
## 99997:      199996  71 4.262680
## 99998:      199997  73 4.290459
## 99999:      199998  60 4.094345
## 100000:      199999  61 4.110874
```

## Evaluating expressions by groups

- If we specify the by= argument, the j argument will be evaluated within each group specified
- This is an especially fast way to evaluate expressions by group levels
- The by= variable is automatically returned as part of the expression

```
## Average age by different gender levels
```

```
cohort[,mean(age),by=gender]
```

```
##      gender      V1
## 1:   Male 64.91831
## 2: Female 65.06756
```

```
## Use a named list to name the variable
```

```
cohort[,list(avg_age=mean(age)),
       by=gender]
```

```
##      gender avg_age
## 1:   Male 64.91831
## 2: Female 65.06756
```

```
## Return multiple columns
cohort[,list(avg_age=mean(age),
             sd_age=sd(age)),
       by=gender]

##   gender avg_age sd_age
## 1:  Male 64.91831 10.007413
## 2: Female 65.06756  9.990478
```

## Evaluating expressions by groups

Question: For the rows after the year 2000, calculate the average and total rbi's for each team

Question: For only the rows after the year 2000, calculate the average rbi's for each league variable (lg)

- Answer

```
baseball_dt[year > 2000,
             list(avg_rbi=mean(rbi,na.rm=T),
                  tot_rbi=sum(rbi,na.rm=T)),
            by=team]

##   team  avg_rbi tot_rbi
## 1:  CLE 31.435897   1226
## 2:  CHA 35.216216   1303
## 3:  BAL 30.270270   1120
## 4:  SEA 26.783333   1607
## 5:  BOS 18.507463   1240
## 6:  TEX 25.697674   1105
## 7:  TBA 27.555556    496
## 8:  NYA 20.189873   1595
## 9:  TOR 44.785714   1254
## 10: DET 31.150000    623
## 11: OAK 19.470588    331
## 12: KCA 17.741935    550
## 13: MIN  8.730769    227
## 14: ANA  1.727273     19
## 15: HOU 34.094340   1807
## 16: NYN 18.658537   1530
## 17: LAN 25.589286   1433
## 18: COL 25.222222   1135
## 19: ARI 26.507692   1723
## 20: SLN 27.711111   1247
```

```
## 21: SFN 35.312500    2260
## 22: ATL 22.104167    1061
## 23: MIL 27.384615     712
## 24: PHI 14.318182     630
## 25: MON 22.777778     205
## 26: CHN 31.851852    1720
## 27: PIT 16.818182     370
## 28: FLO 21.193548     657
## 29: SDN 19.044776    1276
## 30: CIN 22.000000     770
## 31: WAS 13.222222     119
## 32: LAA 27.000000      54
##      team  avg_rbi tot_rbi
```

```
baseball_dt[year > 2000,
             list(avg_rbi=mean(rbi,na.rm=T),
                  tot_rbi=sum(rbi,na.rm=T)),
             by=lg]
```

```
##      lg  avg_rbi tot_rbi
## 1: AL 24.75728   12750
## 2: NL 24.70861   18655
```

## Evaluating expressions by groups

- We may want to group by more than one level
- In that case, we can pass in either a list of column names as objects, or a character vector of column names
- If we pass in a list, we can create a new expression level in the list

*## multiple by variables using list*

```
cohort[,list(age=mean(age)),
        by=list(gender,SES)]
```

```
##      gender  SES      age
## 1:   Male  Med 64.83112
## 2:   Male  Low 64.95914
## 3:   Male High 65.04295
## 4: Female  Med 64.92555
## 5: Female High 65.23148
## 6: Female  Low 65.16029
```

*## multiple by variables using character vector*

```
cohort[,list(age=mean(age)),  
        by=c("gender", "SES")]
```

```
##   gender  SES      age  
## 1:   Male  Med 64.83112  
## 2:   Male  Low 64.95914  
## 3:   Male High 65.04295  
## 4: Female  Med 64.92555  
## 5: Female High 65.23148  
## 6: Female  Low 65.16029
```

*## Use an expression for a by level*

```
cohort[,list(age=mean(age)),  
        by=list(gender, young=age<50)]
```

```
##   gender young      age  
## 1:   Male FALSE 66.22100  
## 2: Female FALSE 66.33298  
## 3:   Male  TRUE 45.22405  
## 4: Female  TRUE 44.99491
```

## Sorting a Data.table

- We often wish to sort the results of an expression such as the one we just saw
- We can use `setorder()` command to sort a data.table in place
- Rather than set the order in place, we can also return the data.table in a new order by calling the `order()` command, then passing it into the `i=` argument.

```
res1 <-  
  baseball_dt[year > 2000,  
              list(avg_rbi=mean(rbi, na.rm=T),  
                  tot_rbi=sum(rbi, na.rm=T)),  
              by=team]
```

```
setorder(res1, avg_rbi)
```

```
res1
```

```
##   team  avg_rbi tot_rbi  
## 1:  ANA  1.727273     19  
## 2:  MIN  8.730769    227  
## 3:  WAS 13.222222    119  
## 4:  PHI 14.318182    630
```

```
## 5: PIT 16.818182    370
## 6: KCA 17.741935    550
## 7: BOS 18.507463   1240
## 8: NYN 18.658537   1530
## 9: SDN 19.044776   1276
## 10: OAK 19.470588    331
## 11: NYA 20.189873   1595
## 12: FLO 21.193548    657
## 13: CIN 22.000000    770
## 14: ATL 22.104167   1061
## 15: MON 22.777778    205
## 16: COL 25.222222   1135
## 17: LAN 25.589286   1433
## 18: TEX 25.697674   1105
## 19: ARI 26.507692   1723
## 20: SEA 26.783333   1607
## 21: LAA 27.000000     54
## 22: MIL 27.384615    712
## 23: TBA 27.555556    496
## 24: SLN 27.711111   1247
## 25: BAL 30.270270   1120
## 26: DET 31.150000    623
## 27: CLE 31.435897   1226
## 28: CHN 31.851852   1720
## 29: HOU 34.094340   1807
## 30: CHA 35.216216   1303
## 31: SFN 35.312500   2260
## 32: TOR 44.785714   1254
##      team   avg_rbi tot_rbi
```

*## Same thing, but no intermediate object*

*## Does not order in place*

```
res1 <-
  baseball_dt[year > 2000,
    list(avg_rbi=mean(rbi,na.rm=T),
         tot_rbi=sum(rbi,na.rm=T)),
    by=team][order(avg_rbi)]
```

## Regression by Group

- Any valid R functions can be evaluated by group....this is powerful!

- We can fit a multivariate regression in the `j` expression and extract the p-values associated with age and mortality
- The first step is to figure out how to get the result you want, then add the `by=` argument in to resolve the expression by group
- In this case, we are fitting a linear regression, then extracting the p-values associated with the mortality predictor

```
## extract the p-values for age and mortality
cohort[,summary(lm(age ~ gender + mortality))$coefficients[2,4],
        by=hosp_id]
```

```
##      hosp_id      V1
## 1:      77 0.121973347
## 2:      45 0.210289025
## 3:       5 0.980924814
## 4:      13 0.887124987
## 5:      70 0.858322923
## 6:      27 0.951418487
## 7:      29 0.491992103
## 8:      85 0.239406643
## 9:      41 0.797902663
## 10:     64 0.603953309
## 11:     43 0.042534060
## 12:     62 0.288284257
## 13:     54 0.795493702
## 14:     69 0.567321025
## 15:     20 0.645297833
## 16:     46 0.282262549
## 17:     19 0.457915985
## 18:      1 0.241606281
## 19:     81 0.567469864
## 20:     55 0.630498498
## 21:     88 0.054499752
## 22:     34 0.064525949
## 23:     15 0.086299501
## 24:     35 0.410806608
## 25:     94 0.557116047
## 26:      2 0.519355657
## 27:     76 0.589302422
## 28:     75 0.206815597
## 29:     66 0.583151953
```



##	30:	91	0.111223202
##	31:	63	0.502365243
##	32:	7	0.550516319
##	33:	22	0.285180780
##	34:	79	0.336277925
##	35:	90	0.951291640
##	36:	95	0.541878016
##	37:	97	0.963290970
##	38:	6	0.780778184
##	39:	36	0.949266143
##	40:	61	0.250958368
##	41:	82	0.566796959
##	42:	25	0.838208685
##	43:	4	0.231794037
##	44:	50	0.333371025
##	45:	33	0.245258738
##	46:	26	0.405456544
##	47:	52	0.961953487
##	48:	12	0.975255867
##	49:	14	0.835948923
##	50:	18	0.971313166
##	51:	31	0.569980012
##	52:	53	0.646286053
##	53:	78	0.886059765
##	54:	59	0.624216847
##	55:	28	0.252300795
##	56:	93	0.020239762
##	57:	87	0.373143181
##	58:	72	0.232146768
##	59:	60	0.298261161
##	60:	51	0.794773690
##	61:	17	0.722502499
##	62:	9	0.080573683
##	63:	73	0.722582454
##	64:	39	0.330502188
##	65:	8	0.518673141
##	66:	100	0.814276270
##	67:	67	0.222394088
##	68:	56	0.011362725
##	69:	48	0.257889098
##	70:	40	0.051644566
##	71:	98	0.331692019
##	72:	57	0.860790323

```
## 73:      80 0.681973253
## 74:       3 0.091395810
## 75:      58 0.012353405
## 76:      74 0.260627155
## 77:      44 0.299418585
## 78:      30 0.432579802
## 79:      92 0.680032805
## 80:      96 0.032909438
## 81:      47 0.573595675
## 82:      84 0.721069120
## 83:      89 0.080750311
## 84:      24 0.118326322
## 85:      49 0.277934332
## 86:      99 0.126635305
## 87:      21 0.812356905
## 88:      83 0.351053376
## 89:      23 0.688500971
## 90:      86 0.926979107
## 91:      42 0.816799942
## 92:      16 0.809076281
## 93:      38 0.158349413
## 94:      10 0.027490918
## 95:      11 0.467260926
## 96:      32 0.241127947
## 97:      37 0.953548590
## 98:      71 0.558666806
## 99:      68 0.009686422
## 100:     65 0.055406935
##      hosp_id      V1
```

## Updating Values inside a Data.Table

- What if we want to add a column into a data.table or perhaps change an existing column?
- We use the := operator in the j argument to update columns inside the data.table
- There are a few different ways to use this operator (next slides will show them)

```
## Create a new variable and save it to the DT
cohort[,l_age:=log(age)]
cohort
```

```
##      patient_id hosp_id age gender mortality SES    l_age
##    1:      100000      77  70   Male        Yes  Med  4.248495
##    2:      100001      45  64   Male        No   Low  4.158883
##    3:      100002       5  76   Male        No   Med  4.330733
##    4:      100003      13  51   Male        No  High  3.931826
##    5:      100004      70  76   Male        No   Low  4.330733
##    ---
## 99996:      199995      60  66   Male        No   Med  4.189655
## 99997:      199996      14  71   Male        No   Low  4.262680
## 99998:      199997      39  73   Male        Yes  Med  4.290459
## 99999:      199998      21  60 Female        No  High  4.094345
## 100000:      199999      81  61   Male        No   Low  4.110874
```

*## to erase the variable we assign it a NULL value*

```
cohort[,l_age:=NULL]
```

```
cohort
```

```
##      patient_id hosp_id age gender mortality SES
##    1:      100000      77  70   Male        Yes  Med
##    2:      100001      45  64   Male        No   Low
##    3:      100002       5  76   Male        No   Med
##    4:      100003      13  51   Male        No  High
##    5:      100004      70  76   Male        No   Low
##    ---
## 99996:      199995      60  66   Male        No   Med
## 99997:      199996      14  71   Male        No   Low
## 99998:      199997      39  73   Male        Yes  Med
## 99999:      199998      21  60 Female        No  High
## 100000:      199999      81  61   Male        No   Low
```

```
names(cohort)
```

```
## [1] "patient_id" "hosp_id"    "age"         "gender"      "mortality"
```

```
## [6] "SES"
```

## Update Multiple Columns

- If we want to update multiple columns, we change the syntax slightly
- Note that any existing column we reference is over-written, and any new column referenced will be created

```
## I use an if-else statement here to create the variable
cohort[, 'l_age:=log(age),
        age_cat=ifelse(age<55, "Young", "Less Young"))]
```

```
cohort
```

```
##      patient_id hosp_id age gender mortality  SES    l_age
age_cat
##      1:      100000      77  70   Male         Yes  Med  4.248495 Less
Young
##      2:      100001      45  64   Male         No   Low  4.158883 Less
Young
##      3:      100002       5  76   Male         No   Med  4.330733 Less
Young
##      4:      100003      13  51   Male         No  High  3.931826
Young
##      5:      100004      70  76   Male         No   Low  4.330733 Less
Young
##      ---

## 99996:      199995      60  66   Male         No   Med  4.189655 Less
Young
## 99997:      199996      14  71   Male         No   Low  4.262680 Less
Young
## 99998:      199997      39  73   Male         Yes   Med  4.290459 Less
Young
## 99999:      199998      21  60 Female         No  High  4.094345 Less
Young
## 100000:      199999      81  61   Male         No   Low  4.110874 Less
Young
```

```
## alternatively, we could have just written two separate lines
```

```
cohort[, l_age:=log(age)]
cohort[, age_cat:=ifelse(age<55, "Young", "Less Young")]
```

## Combining i and j with :=

- If we use the i argument to subset the data then we are only updating the values in that subset when we use :=

```
## I can conditionally update values
```

```
## by combining the first two arguments
```

```
cohort[age > 70 & gender == "Male", risk_cat := "High"]
cohort[!(age > 70 & gender == "Male"), risk_cat := "Low"]
cohort
```

```

##      patient_id hosp_id age gender mortality  SES    l_age
age_cat
##      1:      100000      77  70   Male        Yes  Med  4.248495  Less
Young
##      2:      100001      45  64   Male        No   Low  4.158883  Less
Young
##      3:      100002       5  76   Male        No   Med  4.330733  Less
Young
##      4:      100003      13  51   Male        No  High  3.931826
Young
##      5:      100004      70  76   Male        No   Low  4.330733  Less
Young
##      ---

## 99996:      199995      60  66   Male        No   Med  4.189655  Less
Young
## 99997:      199996      14  71   Male        No   Low  4.262680  Less
Young
## 99998:      199997      39  73   Male        Yes  Med  4.290459  Less
Young
## 99999:      199998      21  60  Female        No  High  4.094345  Less
Young
## 100000:      199999      81  61   Male        No   Low  4.110874  Less
Young
##      risk_cat
##      1:      Low
##      2:      Low
##      3:      High
##      4:      Low
##      5:      High
##      ---
## 99996:      Low
## 99997:      High
## 99998:      High
## 99999:      Low
## 100000:      Low

```

## Group by updates

- We often want to calculate a summary statistic within groups, and then join it back to the original data we calculated it from

- This is very easy by combining the := operator in the j argument with the by= argument
- Here we calculate the average age within hospital, and assign it to the data.table

```
cohort[, avg_age_hosp := mean(age, na.rm=T),
      by=hosp_id]
```

## Exercise: Group by and updates

Question: Using the baseball data, answer this question: Starting after the year 2000, how many players (id denotes players) had an average rbi greater than 40?

Question: Calculate the average rbi value per team and year and save it back to the data.table

- Answer

```
## how many players had average rbi above 40
baseball_dt[year>2000,
             list(avg_rbi=mean(rbi, na.rm=T)),
             by=id][avg_rbi > 40]
```

```
##           id    avg_rbi
## 1:  bellda01  53.57143
## 2:  bicheda01  49.00000
## 3:  boonebr01  80.83333
## 4:  burksel01  48.50000
## 5:  cansejo01  49.00000
## 6:  castivi02  56.50000
## 7:  coninje01  47.40000
## 8:  delgaca01 110.00000
## 9:  evereca01  45.87500
##10:  gonzaju03  52.40000
##11:  justida01  50.00000
##12:  martied01  84.00000
##13:  martiti02  76.40000
##14:  mcgrifr01  50.40000
##15:  olerujo01  60.83333
##16:  oneilpa01  70.00000
##17:  palmera01  97.60000
##18:  ramirma02 114.28571
##19:  ripkeca01  68.00000
##20:  rodriiv01  68.28571
```

```

## 21: thomafr04 70.14286
## 22: thomeji01 101.85714
## 23: valenjo03 54.42857
## 24: vizquom01 50.57143
## 25: willibe02 75.83333
## 26: aloumo01 78.85714
## 27: bagweje01 87.20000
## 28: biggicr01 62.00000
## 29: bondsba01 84.42857
## 30: burnije01 68.14286
## 31: edmonji01 86.42857
## 32: finlest01 52.75000
## 33: floydcl01 55.55556
## 34: gonzalu01 88.14286
## 35: gracema01 47.33333
## 36: greensh01 74.37500
## 37: griffke02 61.57143
## 38: grissma02 60.80000
## 39: jordabr01 43.66667
## 40: kentje01 95.14286
## 41: kleskry01 63.57143
## 42: lopezja01 56.71429
## 43: mcgwima01 64.00000
## 44: nevinph01 50.77778
## 45: piazzmi01 64.85714
## 46: sandere02 63.28571
## 47: sheffga01 94.28571
## 48: snowjt01 40.33333
## 49: sosasa01 98.00000
## 50: stairma01 45.11111
## 51: venturo01 47.40000
## 52: walkela01 67.50000
## 53: walketo04 41.22222
## 54: whitede03 47.00000
## 55: whitero02 47.12500
## 56: willima04 40.33333
## 57: zeileto01 45.20000
##          id    avg_rbi

```

```

## calculate the average rbi by team and year and save it
baseball_dt[,avg_rbi_team_yr:=mean(rbi,na.rm=T),
             by=list(team,year)]

```

## .N variable and Copying

- Data.table is careful to not make extra copies (part of the efficiency)
- If you want to make another copy of a data.table, you should use the `copy()` command explicitly
- Also, there is a special variable called `.N` that can be called in the `j` expression. This is the length of the current group, which can be used to quickly calculate the total rows in a group

```
## Make a copy of a data.table
```

```
cohort2 <- copy(cohort)
```

```
## Use .N
```

```
cohort[, list(avg=mean(age, na.rm=T), .N),  
        by=gender]
```

```
##   gender      avg      N  
## 1:   Male 64.91831 90068  
## 2: Female 65.06756  9932
```

## Setting Keys and Data.table

- We said one reason to use data.table was computational speed
- Data.table has speed gains primarily due to leveraging sorted data to perform binary scans
- This concept is exactly like a telephone book. You can quickly find someone in the telephone book because it is sorted by last name then by first name. That is an example of a binary search. A slower way to do the same search would be to check every name in the phone book to see if it matches the one you are looking for (vector scan)
- In order to leverage a binary search, we must first sort the data.table (just like the phone book must be sorted by last name then first name)



## Setting Keys and Data.table

- We call this sorting “setting the keys” and the keys of the data.table are the variables we sorted it by
- We can sort the table, or “set the keys” with the `setkey()` command
- If we want descending sort order, we need to use the `setorder()` command

```
## Set the key to be hosp_id
```

```
setkey(cohort,hosp_id)
```

```
## notice it is sorted now
```

```
cohort
```

```
##      patient_id hosp_id age gender mortality  SES    l_age
age_cat
##      1:      100017      1  70   Male          No  Med  4.248495 Less
Young
##      2:      100049      1  64   Male        Yes  High  4.158883 Less
Young
##      3:      100094      1  71   Male          No  Med  4.262680 Less
Young
##      4:      100210      1  72   Male          No  Low  4.276666 Less
Young
##      5:      100217      1  75   Male          No  Med  4.317488 Less
Young
##      ---

## 99996:      199342      100  71 Female          No  Med  4.262680 Less
Young
## 99997:      199778      100  57   Male          No  Med  4.043051 Less
Young
## 99998:      199925      100  72   Male          No  Low  4.276666 Less
Young
## 99999:      199952      100  82   Male          No  High 4.406719 Less
Young
## 100000:      199976      100  60 Female          No  Low  4.094345 Less
Young
##      risk_cat avg_age_hosp
##      1:      Low      65.08249
##      2:      Low      65.08249
##      3:      High      65.08249
##      4:      High      65.08249
```

```
##      5:      High      65.08249
##      ---
## 99996:      Low      64.25282
## 99997:      Low      64.25282
## 99998:      High      64.25282
## 99999:      High      64.25282
## 100000:      Low      64.25282
```

*## Set the key to be hosp\_id, then gender*

**setkey**(cohort,hosp\_id,gender)

*## Set the order to be gender, then descending age*

**setorder**(cohort,gender,-age)[ ]

```
##      patient_id hosp_id age gender mortality  SES    l_age
age_cat
##      1:      118862      11  98 Female      No  Low  4.584967 Less
Young
##      2:      173163      17  96 Female      No  Low  4.564348 Less
Young
##      3:      174555      23  96 Female      No  Low  4.564348 Less
Young
##      4:      139345      28  96 Female      No  Low  4.564348 Less
Young
##      5:      128902      56  96 Female      No  Low  4.564348 Less
Young
##      ---
```

```
## 99996:      181651      66  27  Male      No  Low  3.295837
Young
## 99997:      130714      31  26  Male      No  Med  3.258097
Young
## 99998:      166410       4  25  Male      Yes  Med  3.218876
Young
## 99999:      108941      44  23  Male      Yes  High 3.135494
Young
## 100000:      102078      38  22  Male      No  Low  3.091042
Young
```

```
##      risk_cat avg_age_hosp
##      1:      Low      64.75764
##      2:      Low      65.10287
##      3:      Low      64.64581
##      4:      Low      65.27697
##      5:      Low      64.95286
```

```
##      ---
## 99996:      Low      64.87476
## 99997:      Low      64.52484
## 99998:      Low      64.95838
## 99999:      Low      64.96362
## 100000:     Low      65.51812
```

## Leveraging the keys for speed gains

- Once a key has been set in a data.table, we can use that key to make many groupwise operations on the data.table faster
- The first example is subsetting the data.table. Recall we use the `i` argument to subset the data.table. If we pass in a list or a data.table to the `i` argument, it will join to the keys of the data.table effectively subsetting the data.
- an alternative is to use the shortcut function `J()`

```
setkey(cohort,gender)
```

```
## subset down to hosp_id == 11
```

```
## fast way leveraging the keys, all equivalent
```

```
cohort[list("Male")]
```

```
##      patient_id hosp_id age gender mortality  SES    l_age
age_cat
##      1:      126803      41 109   Male          No  Low 4.691348 Less
Young
##      2:      119396      35 108   Male          No  Med 4.682131 Less
Young
##      3:      114797      61 107   Male         Yes  Med 4.672829 Less
Young
##      4:      110129      72 107   Male          No  Med 4.672829 Less
Young
##      5:      156709      54 105   Male          No High 4.653960 Less
Young
##      ---
```

```
## 90064:      181651      66 27   Male          No  Low 3.295837
Young
## 90065:      130714      31 26   Male          No  Med 3.258097
Young
## 90066:      166410       4 25   Male         Yes  Med 3.218876
Young
## 90067:      108941      44 23   Male         Yes High 3.135494
Young
```

```
## 90068:      102078      38  22   Male      No  Low 3.091042
Young
```

```
##      risk_cat avg_age_hosp
```

```
##      1:      High      64.72016
```

```
##      2:      High      65.03490
```

```
##      3:      High      65.21696
```

```
##      4:      High      65.30579
```

```
##      5:      High      64.78392
```

```
##      ---
```

```
## 90064:      Low      64.87476
```

```
## 90065:      Low      64.52484
```

```
## 90066:      Low      64.95838
```

```
## 90067:      Low      64.96362
```

```
## 90068:      Low      65.51812
```

```
cohort[.("Male")]
```

```
##      patient_id hosp_id age gender mortality  SES    l_age
age_cat
```

```
##      1:      126803      41 109   Male      No  Low 4.691348 Less
Young
```

```
##      2:      119396      35 108   Male      No  Med 4.682131 Less
Young
```

```
##      3:      114797      61 107   Male      Yes  Med 4.672829 Less
Young
```

```
##      4:      110129      72 107   Male      No  Med 4.672829 Less
Young
```

```
##      5:      156709      54 105   Male      No  High 4.653960 Less
Young
```

```
##      ---
```

```
## 90064:      181651      66  27   Male      No  Low 3.295837
Young
```

```
## 90065:      130714      31  26   Male      No  Med 3.258097
Young
```

```
## 90066:      166410       4  25   Male      Yes  Med 3.218876
Young
```

```
## 90067:      108941      44  23   Male      Yes  High 3.135494
Young
```

```
## 90068:      102078      38  22   Male      No  Low 3.091042
Young
```

```
##      risk_cat avg_age_hosp
```

```
##      1:      High      64.72016
```

```
##      2:      High      65.03490
##      3:      High      65.21696
##      4:      High      65.30579
##      5:      High      64.78392
##      ---
## 90064:      Low      64.87476
## 90065:      Low      64.52484
## 90066:      Low      64.95838
## 90067:      Low      64.96362
## 90068:      Low      65.51812
```

```
cohort[J("Male")]
```

```
##      patient_id hosp_id age gender mortality  SES    l_age
age_cat
##      1:      126803      41 109   Male          No  Low 4.691348 Less
Young
##      2:      119396      35 108   Male          No  Med 4.682131 Less
Young
##      3:      114797      61 107   Male          Yes  Med 4.672829 Less
Young
##      4:      110129      72 107   Male          No  Med 4.672829 Less
Young
##      5:      156709      54 105   Male          No  High 4.653960 Less
Young
##      ---
```

```
## 90064:      181651      66 27   Male          No  Low 3.295837
Young
## 90065:      130714      31 26   Male          No  Med 3.258097
Young
## 90066:      166410       4 25   Male          Yes  Med 3.218876
Young
## 90067:      108941      44 23   Male          Yes  High 3.135494
Young
## 90068:      102078      38 22   Male          No  Low 3.091042
Young
```

```
##      risk_cat avg_age_hosp
##      1:      High      64.72016
##      2:      High      65.03490
##      3:      High      65.21696
##      4:      High      65.30579
##      5:      High      64.78392
```

```
##      ---
## 90064:      Low      64.87476
## 90065:      Low      64.52484
## 90066:      Low      64.95838
## 90067:      Low      64.96362
## 90068:      Low      65.51812
```

```
## slow way (vector scan)
cohort[gender == "Male"]
```

```
##      patient_id hosp_id age gender mortality  SES      l_age
age_cat
##      1:      126803      41 109   Male          No  Low 4.691348 Less
Young
##      2:      119396      35 108   Male          No  Med 4.682131 Less
Young
##      3:      114797      61 107   Male         Yes  Med 4.672829 Less
Young
##      4:      110129      72 107   Male          No  Med 4.672829 Less
Young
##      5:      156709      54 105   Male          No High 4.653960 Less
Young
##      ---
```

```
## 90064:      181651      66  27   Male          No  Low 3.295837
Young
## 90065:      130714      31  26   Male          No  Med 3.258097
Young
## 90066:      166410       4  25   Male         Yes  Med 3.218876
Young
## 90067:      108941      44  23   Male         Yes High 3.135494
Young
## 90068:      102078      38  22   Male          No  Low 3.091042
Young
```

```
##      risk_cat avg_age_hosp
##      1:      High      64.72016
##      2:      High      65.03490
##      3:      High      65.21696
##      4:      High      65.30579
##      5:      High      64.78392
##      ---
## 90064:      Low      64.87476
## 90065:      Low      64.52484
```

```
## 90066:      Low      64.95838
## 90067:      Low      64.96362
## 90068:      Low      65.51812
```

## Subsetting with multiple keys

- If we have set multiple keys to the data.table, we can pass in multiple arguments to i for the fast subset

```
## subset down to the males with mortality == Yes
```

```
setkey(cohort,gender,mortality)
```

```
cohort[J("Male","Yes")]
```

```
##      patient_id hosp_id age gender mortality  SES    l_age
age_cat
##      1:      114797      61 107   Male        Yes  Med  4.672829 Less
Young
##      2:      191912      73 103   Male        Yes  Med  4.634729 Less
Young
##      3:      108425      64 102   Male        Yes High  4.624973 Less
Young
##      4:      157047      46 101   Male        Yes High  4.615121 Less
Young
##      5:      104316      48 100   Male        Yes  Low  4.605170 Less
Young
##      ---
```

```
## 16452:      117968        1  31   Male        Yes  Med  3.433987
Young
## 16453:      176867      71  31   Male        Yes  Med  3.433987
Young
## 16454:      127488      95  29   Male        Yes  Low  3.367296
Young
## 16455:      166410        4  25   Male        Yes  Med  3.218876
Young
## 16456:      108941      44  23   Male        Yes High  3.135494
Young
```

```
##      risk_cat avg_age_hosp
##      1:      High      65.21696
##      2:      High      65.13546
##      3:      High      64.81524
##      4:      High      64.97073
##      5:      High      65.06635
```

```
##      ---
## 16452:      Low      65.08249
## 16453:      Low      64.31027
## 16454:      Low      64.82656
## 16455:      Low      64.95838
## 16456:      Low      64.96362
```

## Exercise

Question: Using the Baseball data, perform a fast subset (set the key, then use J) of the data only looking at the league NL (lg variable)

- Answer

```
setkey(baseball_dt,lg)
baseball_dt[J("NL")]
```

```
##      id year stint team lg   g  ab  r   h X2b X3b hr rbi
sb cs bb so
##      1: ansonca01 1876      1  CHN NL  66 309 63 110   9   7  2  59
NA NA 12  8
##      2: burdoja01 1876      1  HAR NL  69 309 66  80   9   1  0  23
NA NA 13 16
##      3: forceda01 1876      1  PHN NL  60 284 48  66   6   0  0  17
NA NA  5  3
##      4: forceda01 1876      2  NY3 NL   1   3  0   0   0   0  0  0
NA NA  0  0
##      5: gerhajo01 1876      1  LS1 NL  65 292 33  76  10   3  2  18
NA NA  3  5
##      ---
```

```
## 11374: benitar01 2007      2  FLO NL  34   0  0   0   0   0  0  0
0 0 0 0
## 11375: benitar01 2007      1  SFN NL  19   0  0   0   0   0  0  0
0 0 0 0
## 11376: ausmubr01 2007      1  HOU NL 117 349 38  82  16   3  3  25
6 1 37 74
## 11377: aloumo01 2007      1  NYN NL  87 328 51 112  19   1 13  49
3 0 27 30
## 11378: alomasa02 2007      1  NYN NL   8  22  1   3   1   0  0  0
0 0 0 3
##      ibb hbp sh sf gidp avg_rbi_team_yr
##      1:  NA  NA NA NA  NA      59.33333
##      2:  NA  NA NA NA  NA      31.00000
```



```
##      3:  NA  NA NA NA  NA      24.00000
##      4:  NA  NA NA NA  NA      10.00000
##      5:  NA  NA NA NA  NA      13.50000
##      ---
## 11374:    0    0  0  0    0         0.00000
## 11375:    0    0  0  0    0        34.20000
## 11376:    3    6  4  1   11        19.25000
## 11377:    5    2  0  3   13        20.25000
## 11378:    0    0  0  0    0        20.25000
```

## Group by calculations on the key

- Doing group by calculations on a data.table when the group is the key is faster.
- Thus for very large operations, you may set the key first, then do the group by calculations on the data.table.
- This requires doing a join by passing in a data.table or list to the i argument
- You need to say by=.EACHI to get this behaviour.

```
## mortality
```

```
cohort[,sum(mortality == "Yes")/.N]
```

```
## [1] 0.17989
```

```
## mortality by hospital ID (slower way)
```

```
cohort[,sum(mortality == "Yes")/.N,
        by=hosp_id][1:10]
```

```
##      hosp_id      V1
## 1:         11 0.1832512
## 2:         17 0.1668211
## 3:         23 0.1897074
## 4:         28 0.1584062
## 5:         56 0.2036108
## 6:         73 0.1962151
## 7:         29 0.2007913
## 8:         45 0.1761711
## 9:         60 0.1730580
## 10:        64 0.1961905
```

```
## For an even faster operation, use the key
```

```
## set key
```

```
setkey(cohort,hosp_id)
```

```
## mortality by hosp id, faster version of code
```

```
cohort[J(unique(hosp_id)),
      .(mortality_rate=sum(mortality == "Yes")/.N),
      by=.EACHI][1:10]
```

```
##      hosp_id mortality_rate
##  1:         1      0.1810865
##  2:         2      0.1815320
##  3:         3      0.1776181
##  4:         4      0.2121827
##  5:         5      0.1919087
##  6:         6      0.2141434
##  7:         7      0.1633919
##  8:         8      0.2102161
##  9:         9      0.1836529
## 10:        10      0.1715177
```


## Exercise: Mortality Rates

Question: Using the cohort data, calculate the mortality rates by hospital id. Set the key of the table to hospital id before doing the group by calculation. The mortality rate should be the sum of the mortality equal to Yes values divided by the number of rows for that hospital (use the .N variable for size)

Question: After doing these calculations, create a barplot of this result. Before creating the barplot, sort the result by mortality rate from low to high

- Answer

```
setkey(cohort,hosp_id)
res1 <-
  cohort[,list(mort_rate=sum(mortality == "Yes")/.N),
          by=hosp_id]
setorder(res1,mort_rate)
res1[,barplot(mort_rate)]
```



```
##      [,1]
## [1,]  0.7
## [2,]  1.9
## [3,]  3.1
## [4,]  4.3
## [5,]  5.5
## [6,]  6.7
## [7,]  7.9
## [8,]  9.1
## [9,] 10.3
## [10,] 11.5
## [11,] 12.7
## [12,] 13.9
## [13,] 15.1
## [14,] 16.3
## [15,] 17.5
## [16,] 18.7
## [17,] 19.9
## [18,] 21.1
## [19,] 22.3
## [20,] 23.5
## [21,] 24.7
## [22,] 25.9
```

##	[23,]	27.1
##	[24,]	28.3
##	[25,]	29.5
##	[26,]	30.7
##	[27,]	31.9
##	[28,]	33.1
##	[29,]	34.3
##	[30,]	35.5
##	[31,]	36.7
##	[32,]	37.9
##	[33,]	39.1
##	[34,]	40.3
##	[35,]	41.5
##	[36,]	42.7
##	[37,]	43.9
##	[38,]	45.1
##	[39,]	46.3
##	[40,]	47.5
##	[41,]	48.7
##	[42,]	49.9
##	[43,]	51.1
##	[44,]	52.3
##	[45,]	53.5
##	[46,]	54.7
##	[47,]	55.9
##	[48,]	57.1
##	[49,]	58.3
##	[50,]	59.5
##	[51,]	60.7
##	[52,]	61.9
##	[53,]	63.1
##	[54,]	64.3
##	[55,]	65.5
##	[56,]	66.7
##	[57,]	67.9
##	[58,]	69.1
##	[59,]	70.3
##	[60,]	71.5
##	[61,]	72.7
##	[62,]	73.9
##	[63,]	75.1
##	[64,]	76.3
##	[65,]	77.5

```
## [66,] 78.7
## [67,] 79.9
## [68,] 81.1
## [69,] 82.3
## [70,] 83.5
## [71,] 84.7
## [72,] 85.9
## [73,] 87.1
## [74,] 88.3
## [75,] 89.5
## [76,] 90.7
## [77,] 91.9
## [78,] 93.1
## [79,] 94.3
## [80,] 95.5
## [81,] 96.7
## [82,] 97.9
## [83,] 99.1
## [84,] 100.3
## [85,] 101.5
## [86,] 102.7
## [87,] 103.9
## [88,] 105.1
## [89,] 106.3
## [90,] 107.5
## [91,] 108.7
## [92,] 109.9
## [93,] 111.1
## [94,] 112.3
## [95,] 113.5
## [96,] 114.7
## [97,] 115.9
## [98,] 117.1
## [99,] 118.3
## [100,] 119.5
```

## Data.table and the i argument

- When we pass a list or data.table or J into the i argument, we are technically performing a join which results in subset (like in the previous slides)

- We can specify which of the joined rows are returned with the `mult=` argument
- The options for `mult=` include “all” (the default), “first” and “last”

```
## sort by hospital id then age
```

```
setkey(cohort,hosp_id,age)
```

```
## join all the rows in hospitals one to three
```

```
cohort[J(c(1,2,3))]
```

```
##      patient_id hosp_id age gender mortality  SES    l_age
age_cat risk_cat
##    1:      159292      1  31   Male          No  Med  3.433987
Young      Low
##    2:      117968      1  31   Male          Yes  Med  3.433987
Young      Low
##    3:      139877      1  34   Male          No   Med  3.526361
Young      Low
##    4:      112040      1  36 Female          No   Med  3.583519
Young      Low
##    5:      135266      1  36   Male          No   Med  3.583519
Young      Low
##    ---

## 2917:      113069      3  92   Male          No   Med  4.521789 Less
Young      High
## 2918:      190428      3  92   Male          No   Med  4.521789 Less
Young      High
## 2919:      172256      3  95   Male          No  High  4.553877 Less
Young      High
## 2920:      185063      3  95   Male          No   Med  4.553877 Less
Young      High
## 2921:      192823      3  95   Male          No  High  4.553877 Less
Young      High
##      avg_age_hosp
##    1:      65.08249
##    2:      65.08249
##    3:      65.08249
##    4:      65.08249
##    5:      65.08249
##    ---
## 2917:      65.24230
## 2918:      65.24230
```

```
## 2919:      65.24230
## 2920:      65.24230
## 2921:      65.24230
```

```
## join only to the first row
## youngest person per hospital
cohort[J(c(1,2,3)),mult="first"]
```

```
##   patient_id hosp_id age gender mortality  SES    l_age age_cat
risk_cat
## 1:      159292      1  31   Male          No  Med  3.433987  Young
Low
## 2:      158174      2  33   Male          No  Low  3.496508  Young
Low
## 3:      103987      3  34   Male          No  High 3.526361  Young
Low
##   avg_age_hosp
## 1:      65.08249
## 2:      65.09864
## 3:      65.24230
```

```
## join only to the last row
## oldest person per hospital
cohort[J(c(1,2,3)),mult="last"]
```

```
##   patient_id hosp_id age gender mortality  SES    l_age    age_cat
risk_cat
## 1:      175169      1  98   Male          No  Med  4.584967 Less Young
High
## 2:      100971      2 100   Male          No  Low  4.605170 Less Young
High
## 3:      192823      3  95   Male          No  High 4.553877 Less Young
High
##   avg_age_hosp
## 1:      65.08249
## 2:      65.09864
## 3:      65.24230
```

## Joins in Data.table

- We have already seen an example of a join in data.table when leveraging the keys and using J in the i argument

- We can join two data.tables together using the i argument, and the keys dictate which columns are joined

```
## read in hospital traits for example joins
hosp <- fread("hosp_demo.csv")
## set keys for both tables (these are the things we will merge on)
setkey(hosp,hosp_id)
setkey(cohort,hosp_id)
```

```
## right join: all the hospitals and whatever patients match
cohort[hosp]
```

```
##      patient_id hosp_id age gender mortality SES      l_age
age_cat
##      1:      159292      1  31   Male          No Med 3.433987
Young
##      2:      117968      1  31   Male          Yes Med 3.433987
Young
##      3:      139877      1  34   Male          No Med 3.526361
Young
##      4:      112040      1  36 Female          No Med 3.583519
Young
##      5:      135266      1  36   Male          No Med 3.583519
Young
##      ---

## 99996:      132099      100  89   Male          No Med 4.488636 Less
Young
## 99997:      105295      100  90   Male          No Low 4.499810 Less
Young
## 99998:      184176      100  91   Male          No Low 4.510860 Less
Young
## 99999:      192159      100  94   Male          No Med 4.543295 Less
Young
## 100000:      143580      100  98   Male          Yes Low 4.584967 Less
Young
##      risk_cat avg_age_hosp academic
##      1:      Low      65.08249      Yes
##      2:      Low      65.08249      Yes
##      3:      Low      65.08249      Yes
##      4:      Low      65.08249      Yes
##      5:      Low      65.08249      Yes
##      ---
## 99996:      High      64.25282      No
```



```
## 99997:      High      64.25282      No
## 99998:      High      64.25282      No
## 99999:      High      64.25282      No
## 100000:     High      64.25282      No
```

*## left join: all the patients and whatever hospitals match*  
hosp[cohort]

```
##      hosp_id academic patient_id age gender mortality SES
l_age
##      1:      1      Yes      159292 31  Male      No Med
3.433987
##      2:      1      Yes      117968 31  Male      Yes Med
3.433987
##      3:      1      Yes      139877 34  Male      No Med
3.526361
##      4:      1      Yes      112040 36 Female      No Med
3.583519
##      5:      1      Yes      135266 36  Male      No Med
3.583519
##      ---
```

```
## 99996:      100      No      132099 89  Male      No Med
4.488636
## 99997:      100      No      105295 90  Male      No Low
4.499810
## 99998:      100      No      184176 91  Male      No Low
4.510860
## 99999:      100      No      192159 94  Male      No Med
4.543295
## 100000:      100      No      143580 98  Male      Yes Low
4.584967
```

```
##      age_cat risk_cat avg_age_hosp
##      1:      Young      Low      65.08249
##      2:      Young      Low      65.08249
##      3:      Young      Low      65.08249
##      4:      Young      Low      65.08249
##      5:      Young      Low      65.08249
##      ---
## 99996: Less Young      High      64.25282
## 99997: Less Young      High      64.25282
## 99998: Less Young      High      64.25282
```

```
## 99999: Less Young      High      64.25282
## 100000: Less Young     High      64.25282
```

*## inner join: only the overlapping hospitals and patients*

```
hosp[cohort,nomatch=0]
```

```
##          hosp_id academic patient_id age gender mortality SES
l_age
```

```
##      1:         1      Yes    159292  31   Male          No Med
3.433987
##      2:         1      Yes    117968  31   Male          Yes Med
3.433987
##      3:         1      Yes    139877  34   Male          No Med
3.526361
##      4:         1      Yes    112040  36 Female          No Med
3.583519
##      5:         1      Yes    135266  36   Male          No Med
3.583519
##      ---
```

```
## 99996:      100        No    132099  89   Male          No Med
4.488636
## 99997:      100        No    105295  90   Male          No Low
4.499810
## 99998:      100        No    184176  91   Male          No Low
4.510860
## 99999:      100        No    192159  94   Male          No Med
4.543295
## 100000:      100        No    143580  98   Male          Yes Low
4.584967
```

```
##          age_cat risk_cat avg_age_hosp
```

```
##      1:      Young      Low    65.08249
##      2:      Young      Low    65.08249
##      3:      Young      Low    65.08249
##      4:      Young      Low    65.08249
##      5:      Young      Low    65.08249
##      ---
```

```
## 99996: Less Young      High      64.25282
## 99997: Less Young      High      64.25282
## 99998: Less Young      High      64.25282
## 99999: Less Young      High      64.25282
## 100000: Less Young     High      64.25282
```

*## Full join (all patients and hospitals, even if they don't match)*  
`merge(hosp,cohort,all=TRUE)`

```
##           hosp_id academic patient_id age gender mortality SES
l_age
##      1:         1      Yes    159292  31  Male          No Med
3.433987
##      2:         1      Yes    117968  31  Male          Yes Med
3.433987
##      3:         1      Yes    139877  34  Male          No Med
3.526361
##      4:         1      Yes    112040  36 Female          No Med
3.583519
##      5:         1      Yes    135266  36  Male          No Med
3.583519
##      ---

## 99996:    100        No    132099  89  Male          No Med
4.488636
## 99997:    100        No    105295  90  Male          No Low
4.499810
## 99998:    100        No    184176  91  Male          No Low
4.510860
## 99999:    100        No    192159  94  Male          No Med
4.543295
## 100000:    100        No    143580  98  Male          Yes Low
4.584967
##           age_cat risk_cat avg_age_hosp
##      1:    Young      Low    65.08249
##      2:    Young      Low    65.08249
##      3:    Young      Low    65.08249
##      4:    Young      Low    65.08249
##      5:    Young      Low    65.08249
##      ---
## 99996: Less Young      High    64.25282
## 99997: Less Young      High    64.25282
## 99998: Less Young      High    64.25282
## 99999: Less Young      High    64.25282
## 100000: Less Young      High    64.25282
```

## Conclusion

- Data.table functionality we covered

- Quickly read in files with `fread`
- Use the square brackets to interact with the `data.table`
- The first `i` argument subsets the data or joins if a `data.table` is given
- The second `j` argument accepts any valid `R` expressions
- If `j` resolves to a simple list a `data.table` is returned
- We can update or add columns using the `:=` operator
- We can do any of these operations by group using the `by=` command
- `Data.table` is faster than base `R` in many tasks
- `Data.table` uses memory more efficiently by making less copies of datasets as you work with them
- `Data.table` is one of the fastest ways to merge and work with big data in `R`