

| | Module/framework/package | Name and a brief description of the algorithm | An example of a situation where using the provided GLM implementation provides superior performance compared to that of base R or its equivalent in Python (identify the equivalent in Python) |
|----------------------------------|--------------------------|--|--|
| a. Base R (in the stats library) | stats packages | glm() is used to fit generalized linear models, specifying a symbolic description of the linear predictor and a description of the error distribution. | When working with a dataset exceeding 10 GB, such as customer transaction records, base R's glm() struggles due to memory limitations and single-threaded computation. In contrast, the equivalent in Python—Dask-ML—enables scalable GLM training by distributing the computation across multiple workers, processing data in parallel and in chunks, thus offering significantly better performance. |
| b. Big data version of R | biglm packages, | By processing data in pieces, GLMs using biglm() enable modeling on datasets that don't fit in memory. | When working with large survey or census data that cannot fit in memory, using biglm allows for chunk-wise computation, enabling GLM fitting without crashing or slowing down—something base R's glm() can't handle efficiently. The equivalent in Python would be Dask-ML or PySpark, which also allow for distributed GLM training on large datasets. |
| c. Dask ML | dask-ml | Using Dask arrays/dataframes, parallel/distributed computing implements scalable machine learning including LinearRegression and LogisticRegression. | In scenarios like training a logistic regression model on an over 100GB dataset of IoT sensor data, Dask ML can efficiently process and fit the model using distributed memory, while scikit-learn would crash or require manual data chunking. The equivalent in R is SparkR, which also supports distributed GLMs. |

| | | | |
|-----------------------|--------------|--|--|
| d. Spark R | SparkR | Using distributed computing, <code>spark.glm()</code> fits GLMs on Spark DataFrames, supporting many families such as binomial, gaussian. | <p>For training models on a Spark DataFrame with hundreds of millions of ad impressions, <code>spark.glm()</code> can leverage the distributed environment to train a GLM efficiently. Base R or Python equivalents would be limited by memory or single-machine constraints.</p> <p>The equivalent in Python is <code>pyspark.ml.classification.LogisticRegression</code>.</p> |
| e. Spark optimization | Spark MLlib | Fits GLMs using scalable optimizers such as stochastic gradient descent (SGD) and L-BFGS, designed for large-scale machine learning pipelines. | <p>In a real-time recommendation system where GLMs must be retrained periodically on streaming data from millions of users, Spark MLlib's optimizers like SGD allow fast convergence across a distributed cluster, outperforming single-machine algorithms in Python or R.</p> <p>The equivalent in Python would be scikit-learn, which lacks scalable optimizers like Spark's SGD for distributed data.</p> |
| f. Scikit-learn | Scikit-learn | Implements <code>LogisticRegression()</code> , <code>LinearRegression()</code> , and other GLMs via <code>sklearn.linear_model</code> , with support for various solvers like <code>liblinear</code> , <code>lbfgs</code> , <code>sag</code> , and <code>saga</code> . | <p>When dealing with a clean, moderately sized (~1GB) customer churn dataset, scikit-learn's <code>LogisticRegression</code> can quickly fit a model with multiple solver options, offering better usability and solver flexibility than base R's <code>glm()</code>.</p> <p>The equivalent in R is <code>glm()</code> in the stats package.</p> |